

DATASCI451_Final_Project

2025-04-08

```
library(rstan)
```

```
## Loading required package: StanHeaders
```

```
##
```

```
## rstan version 2.36.0.9000 (Stan version 2.36.0)
```

```
## For execution on a local, multicore CPU with excess RAM we recommend calling
```

```
## options(mc.cores = parallel::detectCores()).
```

```
## To avoid recompilation of unchanged Stan programs, we recommend calling
```

```
## rstan_options(auto_write = TRUE)
```

```
## For within-chain threading using 'reduce_sum()' or 'map_rect()' Stan functions,
```

```
## change 'threads_per_chain' option:
```

```
## rstan_options(threads_per_chain = 1)
```

```
options(mc.cores = 4)
```

```
olympics <- read.csv('01_data_preprocessing/data/olympics.csv')
```

```
head(olympics)
```

```
##   NOC Gold Silver Bronze Total Year X.Athletes X.Events IsHost Country
## 1 AUS   2      0      0      2   1         5         5      0        4
## 2 AUS   2      0      3      5   2         6         6      0        4
## 3 AUS   0      3      1      4   3         6         6      0        4
## 4 AUS   0      2      1      3   7        29        19      0        4
## 5 AUS   3      1      2      6   8        60        31      0        4
## 6 AUS   1      2      1      4   9        29        26      0        4
```

```
# Held out 2024 Data
```

```
olympics_2024 <- olympics[olympics$Year == 33, ]
```

```
head(olympics_2024)
```

```
##   NOC Gold Silver Bronze Total Year X.Athletes X.Events IsHost Country
## 26 AUS  18     19     16    53  33        649     205      0        4
## 37 CHN  40     27     24    91  33        587     211      0        2
## 65 FRA  16     26     22    64  33        801     223      0        5
## 86 JPN  20     12     13    45  33        598     166      0        3
## 113 USA 40     44     42   126  33        854     234      0        1
```

```

olympics <- olympics[olympics$Year != 33, ]

olympics$log_X.Athletes <- log(olympics$X.Athletes + 1)
olympics$log_X.Events <- log(olympics$X.Events + 1)

# Correlation matrix of the variables
cor_matrix <- cor(olympics[, c("X.Athletes", "X.Events", "IsHost")])

library(corrplot)

```

```
## corrplot 0.95 loaded
```

```

pdf("figures/correlation_matrix.pdf", width = 8, height = 6)

corrplot(
  cor_matrix,
  method = "color",
  tl.col = "black",
  tl.srt = 45,
  tl.cex = 0.9,
  tl.offset = 0.8,
  tl.pos = "lt",
  title = "Correlation Matrix of Predictive Variables",
  mar = c(1, 1, 2, 1),
  addgrid.col = "lightgray",
  col = colorRampPalette(c("skyblue", "white", "#004D99"))(100),
  cl.ratio = 0.2,
  cl.align.text = "c"
)
dev.off()

```

```
## pdf
## 2
```

```

pdf("figures/histogram_medals.pdf", width = 8, height = 8)
par(mfrow = c(3, 1), mar = c(4, 4, 2, 1))
hist(olympics$Gold,
  col = "skyblue",
  border = "white",
  main = "Gold Medals",
  xlab = "Medals Counts",
  yaxt = "n",
  ylab = "")
text(x = max(olympics$Gold) * 0.738,
  y = max(hist(olympics$Gold, plot = FALSE)$counts * 0.8),
  labels = paste("Mean =", round(mean(olympics$Gold), 2),
    "\nVar =", round(var(olympics$Gold), 2)),
  col = "red", cex = 1.2, adj = 0)

hist(olympics$Silver,
  col = "skyblue",
  border = "white",

```

```

    main = "Silver Medals",
    xlab = "Medals Counts",
    yaxt = "n",
    ylab = "")
text(x = max(olympics$Silver) * 0.7,
     y = max(hist(olympics$Silver, plot = FALSE)$counts * 0.8),
     labels = paste("Mean =", round(mean(olympics$Silver), 2),
                    "\nVar =", round(var(olympics$Silver), 2)),
     col = "red", cex = 1.2, adj = 0)

hist(olympics$Bronze,
     col = "skyblue",
     border = "white",
     main = "Bronze Medals",
     xlab = "Medals Counts",
     yaxt = "n",
     ylab = "")
text(x = max(olympics$Bronze) * 0.71,
     y = max(hist(olympics$Bronze, plot = FALSE)$counts * 0.8),
     labels = paste("Mean =", round(mean(olympics$Bronze), 2),
                    "\nVar =", round(var(olympics$Bronze), 2)),
     col = "red", cex = 1.2, adj = 0)
dev.off()

```

```

## pdf
## 2

```

$$\begin{aligned}
 \text{Total}_{c,y} &\sim \text{NegBinomial}(\lambda_{c,y}, \phi) \\
 \log(\lambda_{c,y}) &= \alpha_c + \beta_0 + \beta_1 \text{X.Athletes}_{c,y} + \beta_2 \text{X.Events}_{c,y} + \beta_3 \text{IsHost}_{c,y} \\
 \alpha_c &\sim N(0, 5) \\
 \beta_i &\sim N(0, 5) \text{ for } i = 0, 1, 2, 3, 4 \\
 \phi &\sim \text{Gamma}(2, 0.1)
 \end{aligned}$$

```

stan_model_code <- "
data {
  int<lower=0> N; // number of observations
  int<lower=0> C; // number of countries
  int<lower=0> Y; // number of years
  array[C, Y] int<lower=0> Total; // total medals
  array[C, Y] real<lower=0> N_athletes; // number of athletes
  array[C, Y] real<lower=0> N_events; // number of events
  array[C, Y] int<lower=0, upper=1> IsHost; // host country
}
parameters {
  array[C] real alpha; // country-specific intercepts
  array[4] real beta; // coefficients for predictors
  real<lower=0> phi; // dispersion parameter
}
model {
  for (c in 1:C) {

```

```

    for (y in 1:Y) {
      Total[c, y] ~ neg_binomial_2_log(alpha[c] + beta[1] + beta[2] * N_athletes[c, y] + beta[3] * N_events[c, y])
    }
  }

  // Priors
  alpha ~ normal(0, 5);
  phi ~ gamma(2, 0.1);
  for (i in 1:4) {
    beta[i] ~ normal(0, 5);
  }
}
"

```

```

N = nrow(olympics)
C = length(unique(olympics$NOC))
Y = max(olympics$Year)
Gold = matrix(0, nrow = C, ncol = Y)
Silver = matrix(0, nrow = C, ncol = Y)
Bronze = matrix(0, nrow = C, ncol = Y)
N_athletes = matrix(0, nrow = C, ncol = Y)
N_events = matrix(0, nrow = C, ncol = Y)
IsHost = matrix(0, nrow = C, ncol = Y)
for (i in 1:C) {
  country_data = olympics[olympics$Country == i, ]
  for (j in 1:Y) {
    year_data = country_data[country_data$Year == j, ]
    if (nrow(year_data) > 0) {
      Gold[i, j] = year_data$Gold
      Silver[i, j] = year_data$Silver
      Bronze[i, j] = year_data$Bronze
      N_athletes[i, j] = year_data$log_X.Athletes
      N_events[i, j] = year_data$log_X.Events
      IsHost[i, j] = year_data$IsHost
    }
  }
}

gold_data_list = list(
  N = N,
  C = C,
  Y = Y,
  Total = Gold,
  N_athletes = N_athletes,
  N_events = N_events,
  IsHost = IsHost
)

silver_data_list = list(
  N = N,
  C = C,
  Y = Y,

```

```

    Total = Silver,
    N_athletes = N_athletes,
    N_events = N_events,
    IsHost = IsHost
  )

```

```

bronze_data_list = list(
  N = N,
  C = C,
  Y = Y,
  Total = Bronze,
  N_athletes = N_athletes,
  N_events = N_events,
  IsHost = IsHost
)

```

```

model = stan_model(model_code = stan_model_code)

```

```

gold_fit <- sampling(object = model,
  data = gold_data_list,
  iter = 10000,
  warmup = 2000,
  seed = 451)

```

```

silver_fit <- sampling(object = model,
  data = silver_data_list,
  iter = 10000,
  warmup = 2000,
  seed = 451)

```

```

bronze_fit <- sampling(object = model,
  data = bronze_data_list,
  iter = 10000,
  warmup = 2000,
  seed = 451)

```

```

## Warning: There were 15 transitions after warmup that exceeded the maximum treedepth. Increase max_tr
## https://mc-stan.org/misc/warnings.html#maximum-treedepth-exceeded

```

```

## Warning: Examine the pairs() plot to diagnose sampling problems

```

```

# Analysis of the Gold Model
pdf("figures/gold_model.pdf", width = 8, height = 6)
plot(gold_fit)

```

```

## ci_level: 0.8 (80% intervals)

```

```

## outer_level: 0.95 (95% intervals)

```

```
dev.off()
```

```
## pdf
## 2
```

```
pdf("figures/gold_model_trace.pdf", width = 8, height = 6)
stan_trace(gold_fit, nrow = 3, ncol = 4)
dev.off()
```

```
## pdf
## 2
```

```
pdf("figures/gold_model_hist.pdf", width = 8, height = 6)
stan_hist(gold_fit)
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

```
dev.off()
```

```
## pdf
## 2
```

```
summary(gold_fit)$summary
```

##		mean	se_mean	sd	2.5%	25%
##	alpha[1]	4.417132e-01	0.0266474722	2.0312162	-3.57129046	-0.8998364
##	alpha[2]	1.116215e-01	0.0266001053	2.0315221	-3.91026463	-1.2257557
##	alpha[3]	-7.621047e-01	0.0266344804	2.0320469	-4.77660529	-2.1036917
##	alpha[4]	-8.649524e-01	0.0266275307	2.0309427	-4.87306367	-2.2038460
##	alpha[5]	-8.827668e-01	0.0266258715	2.0316342	-4.89068923	-2.2197583
##	beta[1]	-1.980985e+00	0.0266578386	2.0449556	-5.99327607	-3.3325836
##	beta[2]	8.240061e-01	0.0012116423	0.1352333	0.55903164	0.7330318
##	beta[3]	-8.298702e-03	0.0013336193	0.1508828	-0.30494791	-0.1082781
##	beta[4]	3.058219e-01	0.0009001547	0.1190484	0.07363778	0.2261483
##	phi	1.605423e+01	0.0410980947	5.4119446	8.38420072	12.2217337
##	lp__	-3.230419e+02	0.0228407143	2.3153946	-328.42727773	-324.3654839
##		50%	75%	97.5%	n_eff	Rhat
##	alpha[1]	4.372026e-01	1.79005012	4.4180594	5810.323	1.0007792
##	alpha[2]	1.135166e-01	1.45457046	4.1043523	5832.791	1.0007729
##	alpha[3]	-7.635213e-01	0.58738031	3.2325254	5820.751	1.0007968
##	alpha[4]	-8.676188e-01	0.48065806	3.1151518	5817.462	1.0008014
##	alpha[5]	-8.879349e-01	0.46731865	3.1045263	5822.150	1.0007737
##	beta[1]	-1.974073e+00	-0.63873956	2.0843084	5884.613	1.0008530
##	beta[2]	8.230308e-01	0.91297188	1.0927002	12457.138	1.0000221
##	beta[3]	-7.745153e-03	0.09253793	0.2856607	12800.176	0.9999696
##	beta[4]	3.046984e-01	0.38514276	0.5420399	17490.920	1.0000045
##	phi	1.516191e+01	18.77998188	29.3372129	17340.567	1.0001519
##	lp__	-3.227044e+02	-321.35206883	-319.5297042	10276.156	1.0003389

```
# Analysis of the Silver Model
pdf("figures/silver_model.pdf", width = 8, height = 6)
plot(silver_fit)
```

```
## ci_level: 0.8 (80% intervals)
```

```
## outer_level: 0.95 (95% intervals)
```

```
dev.off()
```

```
## pdf
## 2
```

```
pdf("figures/silver_model_trace.pdf", width = 8, height = 6)
stan_trace(silver_fit, nrow = 3, ncol = 4)
dev.off()
```

```
## pdf
## 2
```

```
pdf("figures/silver_model_hist.pdf", width = 8, height = 6)
stan_hist(silver_fit)
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

```
dev.off()
```

```
## pdf
## 2
```

```
summary(silver_fit)$summary
```

##	mean	se_mean	sd	2.5%	25%
## alpha[1]	0.21643170	0.0265477448	2.0339330	-3.7396161	-1.15619307
## alpha[2]	-0.05082450	0.0265625437	2.0344788	-4.0076937	-1.42909782
## alpha[3]	-0.77738152	0.0265401934	2.0331373	-4.7373074	-2.15503057
## alpha[4]	-0.69951951	0.0265527486	2.0333571	-4.6644180	-2.08405573
## alpha[5]	-0.68447342	0.0265461917	2.0330691	-4.6338251	-2.06760524
## beta[1]	-2.13765888	0.0266056278	2.0478966	-6.1118706	-3.53425290
## beta[2]	0.96110434	0.0010885979	0.1216557	0.7195194	0.87930458
## beta[3]	-0.15187818	0.0011626379	0.1317816	-0.4101343	-0.23967128
## beta[4]	0.08663597	0.0009334605	0.1159357	-0.1422264	0.00856328
## phi	22.11422028	0.0638303422	8.1706493	10.8471363	16.39614398
## lp__	-309.46047705	0.0220797454	2.2720483	-314.7245400	-310.78180171
##	50%	75%	97.5%	n_eff	Rhat
## alpha[1]	0.23322106	1.60639516	4.1613699	5869.729	1.000834
## alpha[2]	-0.04062904	1.34324866	3.9026606	5866.337	1.000851
## alpha[3]	-0.76595959	0.61388289	3.1565584	5868.475	1.000839
## alpha[4]	-0.68432264	0.69456464	3.2519117	5864.194	1.000847

```
## alpha[5]    -0.67482682    0.70536400    3.2462458   5865.430 1.000832
## beta[1]     -2.13303610   -0.75415297    1.8537792   5924.736 1.000816
## beta[2]      0.96118297    1.04337083    1.2008926  12489.057 1.000184
## beta[3]     -0.15312491   -0.06383704    0.1101934  12847.558 1.000202
## beta[4]      0.08628104    0.16433763    0.3160563  15425.615 1.000023
## phi         20.60066478    26.08369107    42.1050653  16385.466 1.000277
## lp__        -309.13105838 -307.79725514 -305.9970707 10588.807 1.000101
```

```
# Analysis of the Bronze Model
```

```
pdf("figures/bronze_model.pdf", width = 8, height = 6)
plot(bronze_fit)
```

```
## ci_level: 0.8 (80% intervals)
```

```
## outer_level: 0.95 (95% intervals)
```

```
dev.off()
```

```
## pdf
## 2
```

```
pdf("figures/bronze_model_trace.pdf", width = 8, height = 6)
stan_trace(bronze_fit, nrow = 3, ncol = 4)
dev.off()
```

```
## pdf
## 2
```

```
pdf("figures/bronze_model_hist.pdf", width = 8, height = 6)
stan_hist(bronze_fit)
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

```
dev.off()
```

```
## pdf
## 2
```

```
summary(bronze_fit)$summary
```

```
##              mean      se_mean      sd      2.5%      25%
## alpha[1]  -0.11396049 0.024260212 2.0281412 -4.1424320 -1.4764136
## alpha[2]  -0.35886712 0.024253484 2.0266685 -4.3931789 -1.7182528
## alpha[3]  -0.71351502 0.024253128 2.0275966 -4.7357349 -2.0766754
## alpha[4]  -0.63142871 0.024280330 2.0273180 -4.6739371 -1.9877145
## alpha[5]  -0.73454896 0.024247001 2.0272016 -4.7627661 -2.0933773
## beta[1]   -2.80163711 0.024388792 2.0501625 -6.7741859 -4.2086843
## beta[2]    1.11460335 0.001128093 0.1208378  0.8731453  1.0345871
## beta[3]   -0.16833498 0.001112039 0.1226876 -0.4058331 -0.2501244
```



```
## beta[4]      -0.05020952 0.000866808 0.1069868 -0.2590715 -0.1214385
## phi         31.58707975 0.089657859 11.8595755 15.0917085 23.1844099
## lp__        -293.73243443 0.022187818 2.2731499 -299.0101911 -295.0329132
##              50%          75%          97.5%      n_eff      Rhat
## alpha[1]    -0.09918133 1.27306728 3.81435511 6988.874 1.000133
## alpha[2]    -0.34375133 1.02834819 3.56113180 6982.600 1.000141
## alpha[3]    -0.69511653 0.67904154 3.21682386 6989.202 1.000125
## alpha[4]    -0.61319786 0.75495317 3.29898205 6971.634 1.000124
## alpha[5]    -0.71887904 0.65159498 3.17628881 6990.011 1.000133
## beta[1]     -2.80895867 -1.42820986 1.27287850 7066.364 1.000115
## beta[2]      1.11459348 1.19578772 1.34993528 11474.022 1.000410
## beta[3]     -0.17020966 -0.08708980 0.07679603 12171.975 1.000279
## beta[4]     -0.05084803 0.01993483 0.16182676 15234.037 1.000020
## phi         29.44861237 37.57043973 60.56322056 17496.918 1.000139
## lp__        -293.38217395 -292.07069309 -290.31235813 10496.076 1.000378
```

```
# Predict the data for 2028 using SLR
```

```
# Not used
```

```
pred_data <- function(data, Country, IsHost) {
  noc_data = data[data$Country == Country, ]
  lm.athletes = lm(X.Athletes ~ Year, data = noc_data)
  summary(lm.athletes)
  lm.events = lm(X.Events ~ Year, data = noc_data)
  summary(lm.events)
  pred = data.frame(
    Year = 34,
    log_X.Athletes = log(predict(lm.athletes, newdata = data.frame(Year = 34))),
    log_X.Events = log(predict(lm.events, newdata = data.frame(Year = 34))),
    IsHost = IsHost,
    Country = Country
  )
  return(pred)
}
```

```
# Predict medals for countries in 2024
```

```
post_pred <- function(fit, new_data) {
  # Extract posterior samples
  posterior_samples <- extract(fit)
  alpha_samples <- posterior_samples$alpha[, 1] # vector of MCMC samples for alpha[c]
  beta_samples <- posterior_samples$beta # matrix: iterations x 4 (or list of vectors)
  phi_samples <- posterior_samples$phi # vector of MCMC samples for phi

  # Calculate log(lambda)
  log_lambda <- alpha_samples +
    beta_samples[,1] +
    beta_samples[,2] * new_data$log_X.Athletes +
    beta_samples[,3] * new_data$log_X.Events +
    beta_samples[,4] * new_data$IsHost

  # Calculate lambda
  lambda_pred <- exp(log_lambda)

  # Generate predictions from Poisson distribution
  set.seed(451) # for reproducibility
```

```

total_pred <- rnbino(n = length(lambda_pred),
                    size = phi_samples,
                    mu = lambda_pred)

return(total_pred)
}

# Predict the medals for 2028
predict_for_country <- function(
  gold_fit,
  silver_fit,
  bronze_fit,
  olympics_data,
  country_id,
  is_host,
  country_name = "COUNTRY",
  figure_root_dir = "figures"
) {

  # 1. Predict the data for 2028
  new_data <- data.frame(
    Year = 33,
    log_X.Athletes = log(olympics_data$X.Athletes[olympics_data$Country == country_id]+1),
    log_X.Events = log(olympics_data$X.Events[olympics_data$Country == country_id]+1),
    IsHost = is_host,
    Country = country_name
  )

  # 2. Predict the medals using the posterior predictive distribution
  gold_pred <- post_pred(gold_fit, new_data)
  silver_pred <- post_pred(silver_fit, new_data)
  bronze_pred <- post_pred(bronze_fit, new_data)

  # Create the directory for saving figures
  dir.create(file.path(figure_root_dir, country_name), showWarnings = FALSE, recursive = TRUE)

  # 3. Plot the posterior predictive distribution
  save_hist <- function(pred_values, medal_type = "Gold") {
    pdf_file <- file.path(figure_root_dir, country_name, paste0(tolower(medal_type), "_predictions.pdf"))
    pdf(pdf_file, width = 8, height = 6)

    plotPost(pred_values, xlab = paste(medal_type, "Medals"),
              main = paste("Predicted", medal_type, "Medals for", country_name, "in 2024"),
              )

    dev.off()
  }

  # 4. Plot for each medal type
  save_hist(gold_pred, medal_type = "Gold")
  save_hist(silver_pred, medal_type = "Silver")
  save_hist(bronze_pred, medal_type = "Bronze")
}

```

```

# 5. Return the summary
out_list <- list(
  country = country_name,
  gold_pred = gold_pred,
  silver_pred = silver_pred,
  bronze_pred = bronze_pred
)

return(invisible(out_list))
}

```

```

# Predict for USA
library(bayesboot)
usa_pred <- predict_for_country(
  gold_fit,
  silver_fit,
  bronze_fit,
  olympics_2024,
  country_id = 1,
  is_host = 0,
  country_name = "USA",
  figure_root_dir = "figures"
)

```

```

# Predict for CHN
chn_pred <- predict_for_country(
  gold_fit,
  silver_fit,
  bronze_fit,
  olympics_2024,
  country_id = 2,
  is_host = 0,
  country_name = "CHN",
  figure_root_dir = "figures"
)

```

```

# Predict for JPN
gbr_pred <- predict_for_country(
  gold_fit,
  silver_fit,
  bronze_fit,
  olympics_2024,
  country_id = 3,
  is_host = 0,
  country_name = "JPN",
  figure_root_dir = "figures"
)

```

```

# Predict for AUS
aus_pred <- predict_for_country(
  gold_fit,
  silver_fit,
  bronze_fit,

```

```
olympics_2024,  
country_id = 4,  
is_host = 0,  
country_name = "AUS",  
figure_root_dir = "figures"  
)  
  
# Predict for FRA  
fra_pred <- predict_for_country(  
  gold_fit,  
  silver_fit,  
  bronze_fit,  
  olympics_2024,  
  country_id = 5,  
  is_host = 1,  
  country_name = "FRA",  
  figure_root_dir = "figures"  
)
```